

random-forest-bank_225612018

September 18, 2023

0.1 1. Introduction

1a. Description The data is related with direct marketing campaigns (phone calls) of a Portuguese banking institution. ##### 1b. Summary The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be ('yes') or not ('no') subscribed.

There are four datasets:

- bank-additional-full.csv with all examples (41188) and 20 inputs, ordered by date (from May 2008 to November 2010), very close to the data analyzed in [Moro et al., 2014]
- bank-additional.csv with 10% of the examples (4119), randomly selected from 1), and 20 inputs.
- bank-full.csv with all examples and 17 inputs, ordered by date (older version of this dataset with less inputs).
- bank.csv with 10% of the examples and 17 inputs, randomly selected from 3 (older version of this dataset with less inputs).

The smallest datasets are provided to test more computationally demanding machine learning algorithms (e.g., SVM). Based on data description, bank-additional-full.csv is chosen since it represent all example and 20 inputs.

1c. Goal The classification goal is to predict if the client will subscribe (yes/no) a term deposit (variable y).

```
[1]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

from imblearn.over_sampling import SMOTENC
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, f1_score, recall_score,
precision_score
from sklearn.model_selection import GridSearchCV

import logging
import json
import os
```

```

import sys
import warnings
from urllib.parse import urlparse
import mlflow
import mlflow.sklearn
from mlflow.models import infer_signature

pd.set_option('display.max_columns', 100)
logging.basicConfig(level=logging.WARN)
logger = logging.getLogger(__name__)

```

0.2 2. Exploratory Data Analysis (EDA)

0.2.1 2a. bank client column description:

- 1 - age (numeric)
- 2 - job : type of job (categorical: 'admin.','blue-collar','entrepreneur','housemaid','management','retired','self-employed','services','student','technician','unemployed','unknown')
- 3 - marital : marital status (categorical: 'divorced','married','single','unknown'; note: 'divorced' means divorced or widowed)
- 4 - education (categorical: 'basic.4y','basic.6y','basic.9y','high.school','illiterate','professional.course','university.degree','unknown')
- 5 - default: has credit in default? (categorical: 'no','yes','unknown')
- 6 - housing: has housing loan? (categorical: 'no','yes','unknown')
- 7 - loan: has personal loan? (categorical: 'no','yes','unknown') ##### related with the last contact of the current campaign:
- 8 - contact: contact communication type (categorical: 'cellular','telephone')
- 9 - month: last contact month of year (categorical: 'jan', 'feb', 'mar', ..., 'nov', 'dec')
- 10 - day_of_week: last contact day of the week (categorical: 'mon','tue','wed','thu','fri')
- 11 - duration: last contact duration, in seconds (numeric). Important note: this attribute highly affects the output target (e.g., if duration=0 then y='no'). Yet, the duration is not known before a call is performed. Also, after the end of the call y is obviously known. Thus, this input should only be included for benchmark purposes and should be discarded if the intention is to have a realistic predictive model. ##### other attributes:
- 12 - campaign: number of contacts performed during this campaign and for this client (numeric, includes last contact)
- 13 - pdays: number of days that passed by after the client was last contacted from a previous campaign (numeric; 999 means client was not previously contacted)
- 14 - previous: number of contacts performed before this campaign and for this client (numeric)
- 15 - poutcome: outcome of the previous marketing campaign (categorical: 'failure','nonexistent','success') ##### social and economic context attributes
- 16 - emp.var.rate: [Employment Variation Rate](#) - quarterly indicator (numeric)
- 17 - cons.price.idx: CPI ([Consumer Price Index](#)) - monthly indicator (numeric)
- 18 - cons.conf.idx: CCI ([Consumer Confidence Index](#)) - monthly indicator (numeric)
- 19 - euribor3m: EURIBOR ([Euro Interbank Offer Rate](#)) 3 month rate - daily indicator (numeric)
- 20 - nr.employed: number of employees - quarterly indicator (numeric) ##### Output variable (desired target):
- 21 - y - has the client subscribed a term deposit? (binary: 'yes','no')

[2]:

```
bank_raw = pd.read_csv('D:/Kuliah/semester_3/kecerdasan_buatan/Pertemuan_2/
↳data-society-bank-marketing-data/bank-additional-full.csv',delimiter=';')
bank_raw.head()
```

```
[2]:
```

	age	job	marital	education	default	housing	loan	contact	\
0	56	housemaid	married	basic.4y	no	no	no	telephone	
1	57	services	married	high.school	unknown	no	no	telephone	
2	37	services	married	high.school	no	yes	no	telephone	
3	40	admin.	married	basic.6y	no	no	no	telephone	
4	56	services	married	high.school	no	no	yes	telephone	

	month	day_of_week	duration	campaign	pdays	previous	poutcome	\
0	may	mon	261	1	999	0	nonexistent	
1	may	mon	149	1	999	0	nonexistent	
2	may	mon	226	1	999	0	nonexistent	
3	may	mon	151	1	999	0	nonexistent	
4	may	mon	307	1	999	0	nonexistent	

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	y
0	1.1	93.994	-36.4	4.857	5191.0	no
1	1.1	93.994	-36.4	4.857	5191.0	no
2	1.1	93.994	-36.4	4.857	5191.0	no
3	1.1	93.994	-36.4	4.857	5191.0	no
4	1.1	93.994	-36.4	4.857	5191.0	no

The table schema is correctly inferred

```
[3]: bank_raw.dtypes
```

```
[3]:
```

age	int64
job	object
marital	object
education	object
default	object
housing	object
loan	object
contact	object
month	object
day_of_week	object
duration	int64
campaign	int64
pdays	int64
previous	int64
poutcome	object
emp.var.rate	float64
cons.price.idx	float64
cons.conf.idx	float64
euribor3m	float64

```
nr.employed    float64
y              object
dtype: object
```

2b. Descriptive Statistics From the statistic descriptive, we knew that data consist of 41188 records and described below:

1. Age averaging in 40 years old, with min 17 years old and max 98 years old.
2. The most frequently job appear is admin with 10422 occurrence.
3. The most frequently marital status appear is married with 24928 occurrence.
4. The most frequently education level appear is university degree with 12168 occurrence.
5. Majority has no credit in default with 32588 occurrence.
6. Majority still has housing loan with 21576 occurrence.
7. Majority has no personal loan with 33950 occurrence.
8. Majority is contacted via cellular with 26144 occurrence.
9. Majority has last contacted in may with 26144 occurrence.
10. Mostly client contacted at Thursday with 8623 occurrence.
11. Duration of calls averaging in 258.29 second, but we wont use this feature since its strongly correlated with y.
12. Client commonly contacted 2 to 3 times within campaign period, with min 1 and max 56.
13. Majority client are not contacted from previous marketing campaign, and client that has been contacted previously having average of 6 days after previous marketing campaign with min 0 and max 27.
14. From previous campaign, client having average of being 1 or 2 times contacted, with min 1 and max 7.
15. From previous campaign, the outcome is most likely failure with 4252 occurrence from total 5625 occurrence.
16. Employee Variation Rate averaging in 0.081886 with min -3.400000 and max 1.400000 (quarterly basis).
17. Consumer Price Index (CPI) averaging in 93.575664 with min 92.201000 and max 94.767000 (monthly basis).
18. Consumer Confidence Index (CCI) averaging in -40.502600 with min -50.800000 max -26.900000 (monthly basis).
19. Euro Interbank Offer Rate (EURIBOR) averaging in 3.621291 with min 0.634000 and max 5.045000 (3 months rate-daily basis).
20. Number of Emploeyess averaging in 5167 with min 4964 and max 5228 (quarterly basis).
21. The outcome has imbalance class with no response accounted for 36548 occurrence.

```
[4]: print(bank_raw.describe(include='all'))
```

	age	job	marital	education	default	housing	\
count	41188.00000	41188	41188	41188	41188	41188	
unique	NaN	12	4	8	3	3	
top	NaN	admin.	married	university.degree	no	yes	
freq	NaN	10422	24928	12168	32588	21576	
mean	40.02406	NaN	NaN	NaN	NaN	NaN	
std	10.42125	NaN	NaN	NaN	NaN	NaN	
min	17.00000	NaN	NaN	NaN	NaN	NaN	

25%	32.00000	NaN	NaN	NaN	NaN	NaN
50%	38.00000	NaN	NaN	NaN	NaN	NaN
75%	47.00000	NaN	NaN	NaN	NaN	NaN
max	98.00000	NaN	NaN	NaN	NaN	NaN

	loan	contact	month	day_of_week	duration	campaign \
count	41188	41188	41188	41188	41188.000000	41188.000000
unique	3	2	10	5	NaN	NaN
top	no	cellular	may	thu	NaN	NaN
freq	33950	26144	13769	8623	NaN	NaN
mean	NaN	NaN	NaN	NaN	258.285010	2.567593
std	NaN	NaN	NaN	NaN	259.279249	2.770014
min	NaN	NaN	NaN	NaN	0.000000	1.000000
25%	NaN	NaN	NaN	NaN	102.000000	1.000000
50%	NaN	NaN	NaN	NaN	180.000000	2.000000
75%	NaN	NaN	NaN	NaN	319.000000	3.000000
max	NaN	NaN	NaN	NaN	4918.000000	56.000000

	pdays	previous	poutcome	emp.var.rate	cons.price.idx \
count	41188.000000	41188.000000	41188	41188.000000	41188.000000
unique	NaN	NaN	3	NaN	NaN
top	NaN	NaN	nonexistent	NaN	NaN
freq	NaN	NaN	35563	NaN	NaN
mean	962.475454	0.172963	NaN	0.081886	93.575664
std	186.910907	0.494901	NaN	1.570960	0.578840
min	0.000000	0.000000	NaN	-3.400000	92.201000
25%	999.000000	0.000000	NaN	-1.800000	93.075000
50%	999.000000	0.000000	NaN	1.100000	93.749000
75%	999.000000	0.000000	NaN	1.400000	93.994000
max	999.000000	7.000000	NaN	1.400000	94.767000

	cons.conf.idx	euribor3m	nr.employed	y
count	41188.000000	41188.000000	41188.000000	41188
unique	NaN	NaN	NaN	2
top	NaN	NaN	NaN	no
freq	NaN	NaN	NaN	36548
mean	-40.502600	3.621291	5167.035911	NaN
std	4.628198	1.734447	72.251528	NaN
min	-50.800000	0.634000	4963.600000	NaN
25%	-42.700000	1.344000	5099.100000	NaN
50%	-41.800000	4.857000	5191.000000	NaN
75%	-36.400000	4.961000	5228.100000	NaN
max	-26.900000	5.045000	5228.100000	NaN

```
[5]: print("how many days after previous marketing campaign")
print(bank_raw[bank_raw['pdays'] != 999]['pdays'].describe())
print()
```

```
print("how about client that has not been contacted before")
print(bank_raw[bank_raw['pdays'] == 999]['pdays'].describe())
```

how many days after previous marketing campaign

```
count    1515.000000
mean      6.014521
std       3.824906
min       0.000000
25%       3.000000
50%       6.000000
75%       7.000000
max       27.000000
```

Name: pdays, dtype: float64

how about client that has not been contacted before

```
count    39673.0
mean     999.0
std       0.0
min     999.0
25%     999.0
50%     999.0
75%     999.0
max     999.0
```

Name: pdays, dtype: float64

```
[6]: print("how many times client is contacted from previous campaign")
print(bank_raw[bank_raw['pdays'] != 999]['previous'].describe())
```

how many times client is contacted from previous campaign

```
count    1515.000000
mean      1.660726
std       0.934306
min       1.000000
25%       1.000000
50%       1.000000
75%       2.000000
max       7.000000
```

Name: previous, dtype: float64

```
[7]: print("from previous campaign, how is the outcome?")
print(bank_raw[bank_raw['poutcome'] != 'nonexistent']['poutcome'].describe())
```

from previous campaign, how is the outcome?

```
count      5625
unique       2
top      failure
freq      4252
```

Name: poutcome, dtype: object

2c. Barchart for Categorical Feature From barchart created above, several point can be derived:

1. 80% client jobs consist of admin, blue-collars, technician, service, and management.
2. 80% client education consist of university degree and high school.
3. Client that has default loan is much fewer and doesnt fall below 80% of the data. Thus it is highly imbalance.
4. Client that has personal loan is much fewer and doesnt fall below 80% of the data. Thus it is highly imbalance.
5. 80% of client were contacted in may-august.
6. Client is uniformly called during the week.

```
[8]: bank_cat = bank_raw.select_dtypes(include=['object'])
bank_cat.head()
```

```
[8]:      job  marital  education  default  housing  loan  contact  month  \
0  housemaid  married   basic.4y      no      no    no  telephone  may
1  services  married  high.school  unknown      no    no  telephone  may
2  services  married  high.school      no    yes    no  telephone  may
3   admin.  married   basic.6y      no      no    no  telephone  may
4  services  married  high.school      no      no   yes  telephone  may

   day_of_week  poutcome  y
0         mon  nonexistent  no
1         mon  nonexistent  no
2         mon  nonexistent  no
3         mon  nonexistent  no
4         mon  nonexistent  no
```

```
[9]: for column in bank_cat:
    plt.figure(figsize=(15,4))

    # Calculate value counts and sort by descending order
    value_counts = bank_cat[column].value_counts().sort_values(ascending=False)

    # Create bar chart
    value_counts.plot(kind='bar', color='blue', alpha=0.5)





    # Calculate cumulative sums and convert to percentage of total
    cumulative_sums = value_counts.cumsum() / value_counts.sum() * 100




    # Create Pareto line
    cumulative_sums.plot(kind='line', marker='o', color='red', secondary_y=True)





    # Add dotted line at 80%
    plt.axhline(y=80, color='k', linestyle='--')

    plt.title(column, fontdict={'fontsize': 20})
```

```
plt.show()
```







2c. Boxplot for Numerical Feature From boxplot created above, several point can be derived:

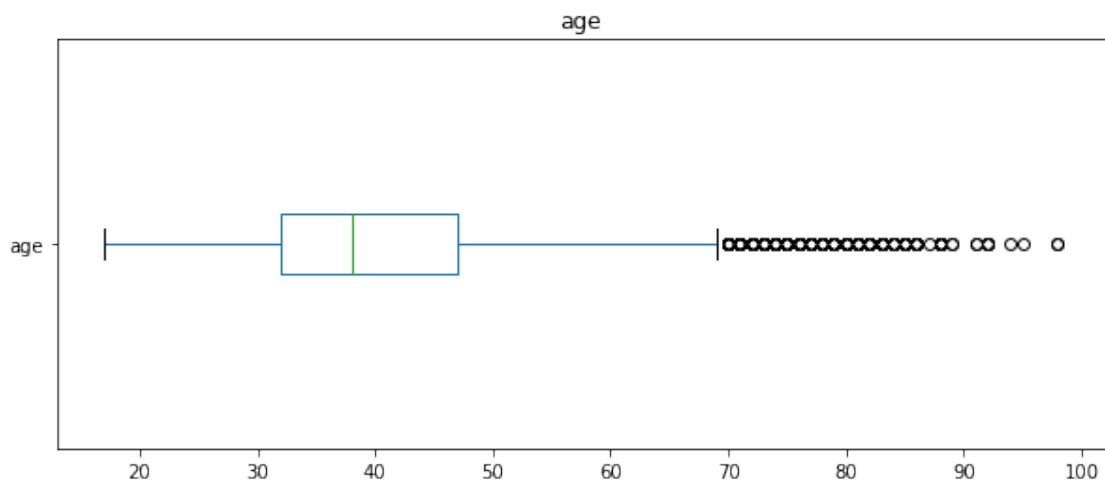
1. Outliers exist in age, duration, campaign, pdays, previous, and CCI


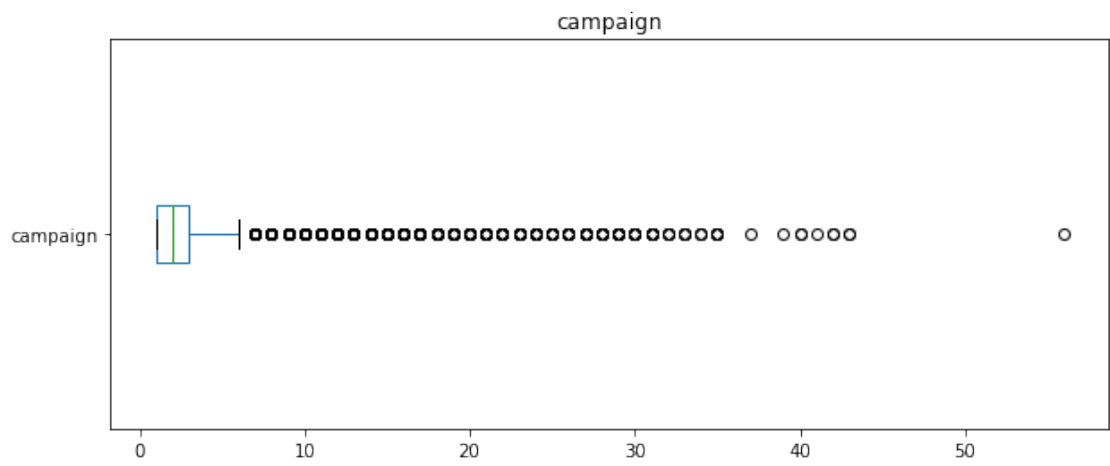

```
[10]: bank_num = bank_raw.select_dtypes(exclude=['object'])
      bank_num.head()
```

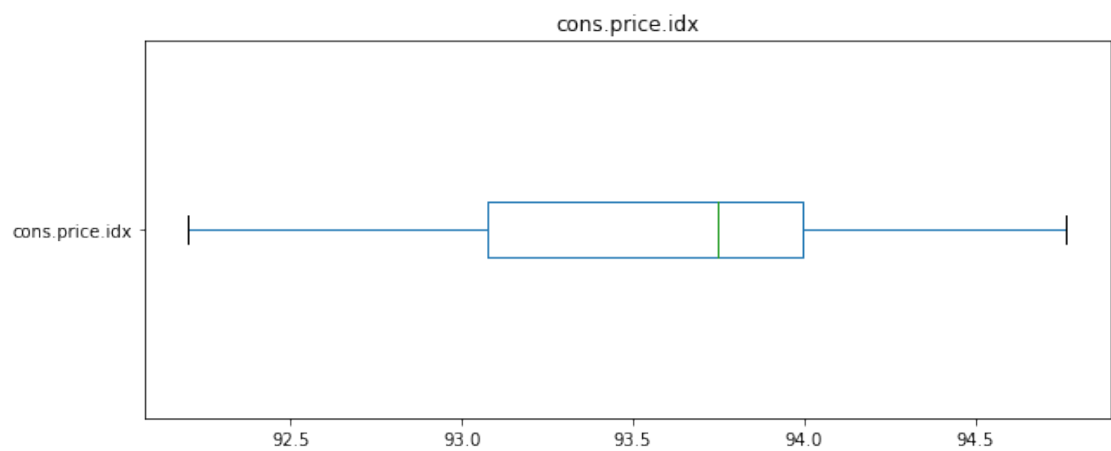
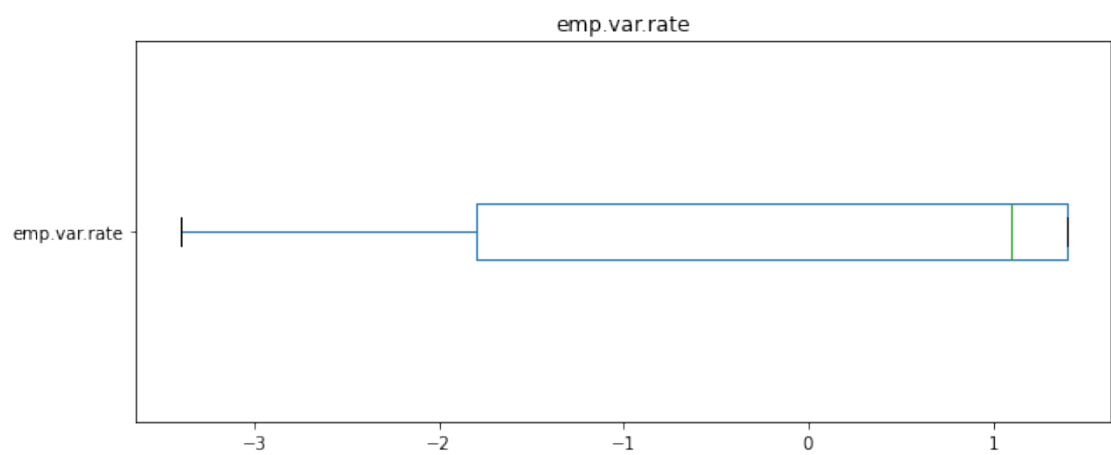

```
[10]:   age  duration  campaign  pdays  previous  emp.var.rate  cons.price.idx  \
0    56        261         1     999         0          1.1         93.994
1    57        149         1     999         0          1.1         93.994
2    37        226         1     999         0          1.1         93.994
3    40        151         1     999         0          1.1         93.994
4    56        307         1     999         0          1.1         93.994

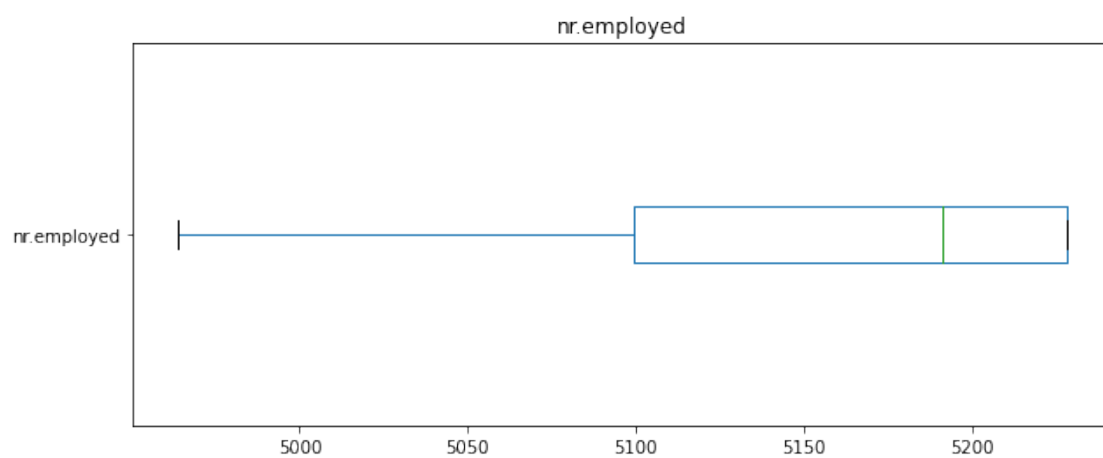
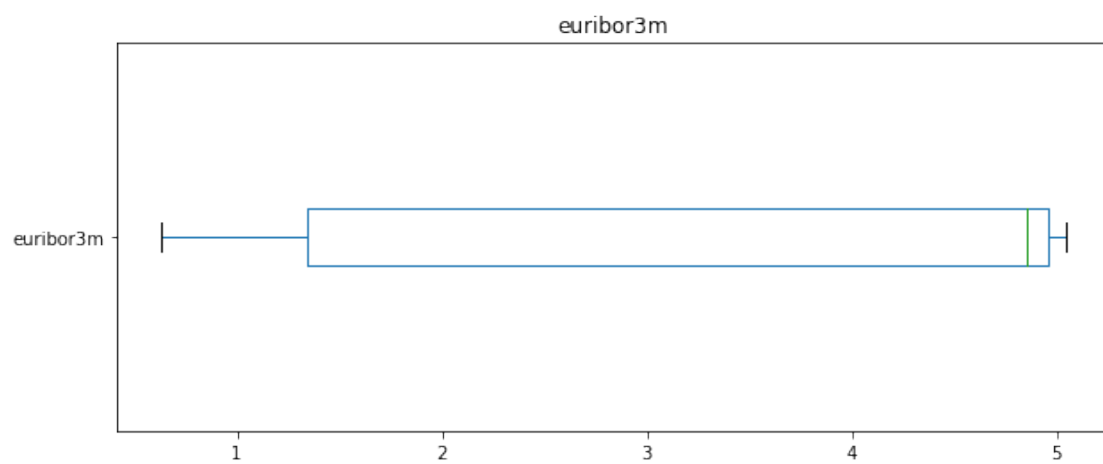
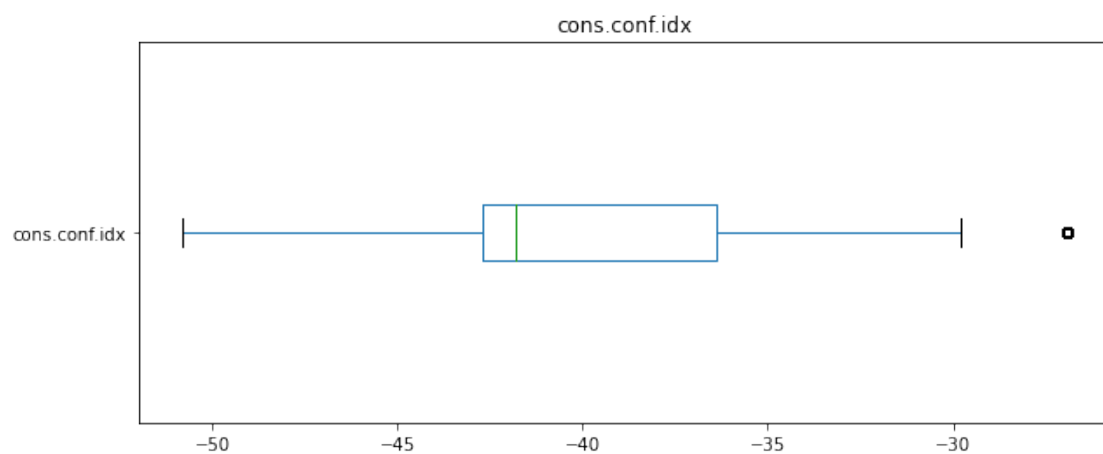
      cons.conf.idx  euribor3m  nr.employed
0          -36.4      4.857      5191.0
1          -36.4      4.857      5191.0
2          -36.4      4.857      5191.0
3          -36.4      4.857      5191.0
4          -36.4      4.857      5191.0
```

```
[11]: for column in bank_num:
      plt.figure(figsize=(10,4))
      bank_num.boxplot([column], vert=False, grid=False)
      plt.title(column)
      plt.show()
```









0.2.2 3. Data Preprocessing

In this script, Random forest classifier is the proposed model, thus before using it, assumption must be satisfied.

1. Sampling is representative.
2. Missing value should be handled from training the model.
3. Data should contain some actual values in feature variables.

Note * There should be keep in mind that a Very unbalanced data in columns may produce bias since most tree may lean towards the biggest data proportion, thus ignoring the others. * No formal distributional assumptions, random forests are non-parametric and can thus handle skewed and multi-modal data as well as categorical data that are ordinal or non-ordinal.

Reference:

<https://stats.stackexchange.com/questions/59124/random-forest-assumptions>

<https://www.simplilearn.com/tutorials/data-science-tutorial/random-forest-in-r>

<https://bccvl.org.au/algorithms-exposed-random-forest/>

3.1. Sampling is Representative Stratified sampling is yet used, it potentially used in `train_test_split` and random forest classifier.

3.2. Check Missing Value There is no missing value found on `bank_raw` data, thus second assumption is satisfied

```
[12]: bank_raw.isna().sum()
```

```
[12]: age          0
      job          0
      marital      0
      education    0
      default      0
      housing      0
      loan         0
      contact      0
      month        0
      day_of_week  0
      duration     0
      campaign     0
      pdays        0
      previous     0
      poutcome     0
      emp.var.rate  0
      cons.price.idx 0
      cons.conf.idx 0
      euribor3m    0
      nr.employed  0
```

```
y          0
dtype: int64
```

3.3. Data Contain Actual Value The use of raw data represent the actual value, thus third assumption is satisfied

3.4. Data Preparation

- Remove duration
- Label Encoding the dataset
- Upsampling minority target
- Split dataset into independent and dependent dataset

```
[13]: bank_rm_duration = bank_raw.drop("duration", axis=1)
```

```
[14]: # Separate features and target
X = bank_rm_duration.drop('y', axis=1)
y = bank_rm_duration['y']

# Create a boolean list indicating which columns are categorical
categorical_features = (X.dtypes == object).tolist()

# Initialize SMOTENC with the correct categorical features
smote_nc = SMOTENC(categorical_features=categorical_features, random_state=0)

# Perform resampling
X_resampled, y_resampled = smote_nc.fit_resample(X, y)
```

```
[15]: # Perform Dummy Variable Encoding
# This method is similar to one-hot encoding but it drops one dummy variable
↳ from the results to avoid multicollinearity, a situation in which two or
↳ more variables are highly correlated.
#bank_enc = pd.get_dummies(bank_rm_duration,
↳ columns=['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'day_of_w
↳ drop_first=True)
bank_enc = bank_rm_duration.
↳ drop(['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'day_of_week
↳ axis=1)
bank_enc['y'] = bank_enc['y'].map({'yes': 1, 'no': 0})

#X_resampled = pd.get_dummies(X_resampled,
↳ columns=['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'day_of_w
↳ drop_first=True)
```



```
X_resampled = X_resampled.  
    ↪drop(['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact', 'month', 'day_of_week',  
    ↪axis=1)  
y_resampled = y_resampled.map({'yes': 1, 'no': 0})
```

```
[16]: X_resampled.head()
```

```
[16]:
```

	age	campaign	pdays	previous	emp.var.rate	cons.price.idx	\
0	56	1	999	0	1.1	93.994	
1	57	1	999	0	1.1	93.994	
2	37	1	999	0	1.1	93.994	
3	40	1	999	0	1.1	93.994	
4	56	1	999	0	1.1	93.994	

	cons.conf.idx	euribor3m	nr.employed
0	-36.4	4.857	5191.0
1	-36.4	4.857	5191.0
2	-36.4	4.857	5191.0
3	-36.4	4.857	5191.0
4	-36.4	4.857	5191.0

```
[37]: #X = bank_enc.drop(["y"], axis=1)  
#y = bank_enc["y"]  
X = X_resampled  
y = y_resampled
```

0.2.3 4. Random Forest Classifier

A random forest classifier.

A random forest is a meta estimator that fits a number of decision tree classifiers on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control overfitting. The sub-sample size is controlled with the `max_samples` parameter if `bootstrap=True` (default), otherwise the whole dataset is used to build each tree.

4.1. Parameter Tunning and Model Training

```
[38]: rand_seed = 123  
np.random.seed(rand_seed)
```

```
[39]: # Assuming X is your feature matrix and y are your labels  
# Generate a random sample for training and testing  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
    ↪random_state=rand_seed)
```

Fitting 3 folds for each of 540 candidates, totalling 1620 fits Best parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 2, 'min_samples_split': 10, 'n_estimators': 200}

```
[20]: ##### Train a Random Forest classifier

##### Define the parameter grid
param_grid = {
    'n_estimators': [100, 200, 300, 400, 500],
    'max_depth': [None, 10, 20, 30, 40, 50],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt']
}

##### Create a base model
rf = RandomForestClassifier(random_state=rand_seed)

##### Instantiate the grid search model
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid,
                           cv=3, n_jobs=-1, verbose=2)

##### Fit the grid search to the data
grid_search.fit(X_train, y_train)

##### Get the best parameters
best_params = grid_search.best_params_

print("Best parameters: ", best_params)
```

Fitting 3 folds for each of 540 candidates, totalling 1620 fits
 Best parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 4, 'min_samples_split': 2, 'n_estimators': 500}

```
[40]: def eval_metrics(actual, pred):
    accuracy = accuracy_score(actual, pred)
    f1 = f1_score(actual, pred)
    recall = recall_score(actual, pred)
    precision = precision_score(actual, pred)
    return accuracy, f1, recall, precision
```

```
[41]: mlflow.set_tracking_uri("http://localhost:5000")
mlflow.set_experiment("Bank_Marketing")
```

```
[41]: <Experiment: artifact_location='mlflow-artifacts:/318929191823953936',
creation_time=1694958964337, experiment_id='318929191823953936',
last_update_time=1694958964337, lifecycle_stage='active', name='Bank_Marketing',
tags={}>
```

```
[42]: if __name__ == "__main__":
    warnings.filterwarnings("ignore")
    rand_seed = 123
```

```

np.random.seed(rand_seed)

# Assuming X is your feature matrix and y are your labels
# Generate a random sample for training and testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=rand_seed)

with mlflow.start_run(run_name="bank_num"):
    clf = RandomForestClassifier(**best_params)
    clf.fit(X_train, y_train)

    # Feature Importance Visualization
    # Assuming clf is your trained model and feature_names is the list of
↪feature names
    feature_importance = clf.feature_importances_
    feature_names = np.array(X.columns)

    # Sort features by importance
    sorted_idx = np.argsort(feature_importance)

    # Create a horizontal bar plot
    plt.figure(figsize=(15,10))
    plt.barh(range(len(sorted_idx)), feature_importance[sorted_idx],
↪align='center')
    plt.yticks(range(len(sorted_idx)), feature_names[sorted_idx])
    plt.xlabel('Importance')
    plt.title('Feature Importance')

    # Save the figure as a PNG
    if not os.path.exists("images"):
        os.mkdir("images")

    plt.savefig("feature_importance.png")
    plt.show()

    # Test Random Forest Classification Model
    y_pred = clf.predict(X_test)

    (accuracy, f1, recall, precision) = eval_metrics(y_test, y_pred)

    print(f"Accuracy: {accuracy}")
    print(f"F1 Score: {f1}")
    print(f"Recall: {recall}")
    print(f"Precision: {precision}")

    mlflow.log_param("accuracy", accuracy)

```


```

mlflow.log_param("f1 score", f1)
mlflow.log_param("recall", recall)
mlflow.log_param("precision", precision)
#mlflow.log_artifact("feature_importance.json")
mlflow.log_artifact("feature_importance.png")

predictions = clf.predict(X_train)
signature = infer_signature(X_train, predictions)
tracking_url_type_store = urlparse(mlflow.get_tracking_uri()).scheme

# Model registry does not work with file store
if tracking_url_type_store != "file":
    mlflow.sklearn.log_model(clf, "model",
    ↪registered_model_name="BankMarketing", signature=signature)
else:
    mlflow.sklearn.log_model(clf, "model", signature=signature)

```



Accuracy: 0.8864569083447332
 F1 Score: 0.8856906762153973
 Recall: 0.8805970149253731
 Precision: 0.8908436071478044

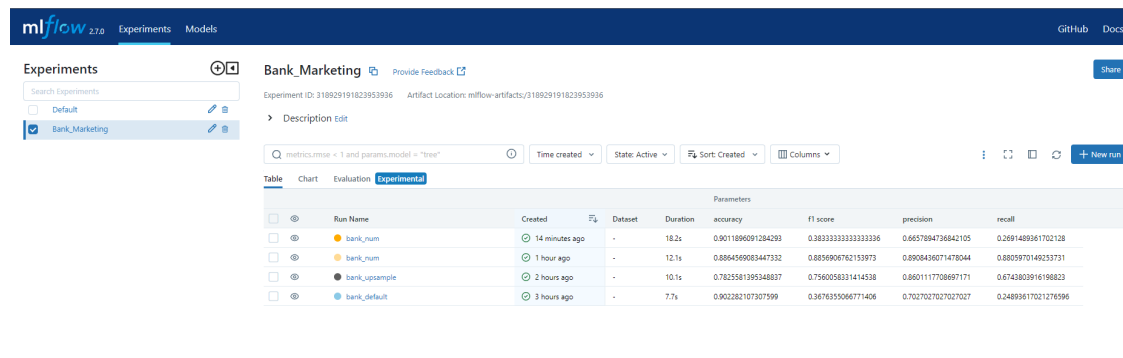
Registered model 'BankMarketing' already exists. Creating a new version of this model...

2023/09/18 00:24:37 INFO mlflow.tracking._model_registry.client: Waiting up to 300 seconds for model version to finish creation. Model name: BankMarketing, version 12

Created version '12' of model 'BankMarketing'.

```
[43]: from IPython.display import Image
      Image("D:
        ↳\Kuliah\semester_3\kecerdasan_buatan\Pertemuan_2\data-society-bank-marketing-data\Model_Com
        ↳png")
```

[43]:



Run Name	Created	Dataset	Duration	accuracy	f1 score	precision	recall
bank_num	14 minutes ago	-	18.2s	0.9011896091284293	0.3833333333333333	0.6657894736842105	0.2691489361702128
bank_num	1 hour ago	-	12.1s	0.8864569083447332	0.8856906762153973	0.8908436071478044	0.8805970149233731
bank_upsample	2 hours ago	-	10.1s	0.7825581395348837	0.7560058331414538	0.8601117708697171	0.6743803916198823
bank_default	3 hours ago	-	7.7s	0.902282107307599	0.3676355066771406	0.7027027027027027	0.24893617021276596

4.2. Model Interpretation Lets learn about how to interpret the metrics

1 - Accuracy is suitable with balanced dataset when there are an equal number of observations in each class which isn't common in real-life problems.

2 - Precision is important when the cost of false positives is high.

3 - Recall is important when the cost of false negatives is high.

4 - F1 score considers both the precision and recall.

Looking back at case at hand, this dataset related with direct marketing campaigns (via phone calls) of a Portuguese banking institution. The goal is to predict if the client will subscribe (yes/no) a term deposit (variable y). Based on this information, it is assumed that in the campaign, client who likely subscribed and not will undergo different approach. Thus, both predictions are equally important.

a. First Training First attempt is performed using dummy variables. Using this set of parameter {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 2, 'min_samples_split': 10, 'n_estimators': 200} yield:

Accuracy: 0.902282107307599

F1 Score: 0.3676355066771406

Recall: 0.24893617021276596

Precision: 0.7027027027027027

The recall is very low, thus affecting F1 score also low. This may happen since the model uses imbalance target, lets try upsampling the minor target (no)

b. Second Training Second attempt is performed using dummy variables and SMOTENC upsampling method. Using this set of parameter {'max_depth': 30, 'max_features': 'auto', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 500} yield:

Accuracy: 0.7825581395348837
F1 Score: 0.7560058331414538
Recall: 0.6743803916198823
Precision: 0.8601117708697171

This result is better than previous attempt since recall value is substantially increased as the result of upsampling using SMOTENC.

c. Third Training Third attempt is performed using only numerical feature on non upsampling data, using this set of parameter {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 4, 'min_samples_split': 2, 'n_estimators': 500} yield:

Accuracy: 0.9011896091284293 F1 Score: 0.38333333333333336 Recall: 0.2691489361702128 Precision: 0.6657894736842105

As expected, using imbalance dataset produce lower recall, thus lowering F1 score which made upsampling method more preferable

d. Forth Training Fourth attempt is performed using only numerical feature on SMOTENC upsampling data, using this set of parameter {'max_depth': 30, 'max_features': 'auto', 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 200} yield:

Accuracy: 0.8864569083447332
F1 Score: 0.8856906762153973
Recall: 0.8805970149253731
Precision: 0.8908436071478044

The performance is considerably better than previous attempt with high F1 Score, therefore this model is able to discriminate outcome (y) better than previous. Thus, this model is selected

Reference:

<https://medium.com/analytics-vidhya/what-precision-recall-f1-score-and-accuracy-can-tell-you-fe1eablada5a>

<https://datascience.stackexchange.com/questions/105089/how-f1-score-is-good-with-unbalanced-dataset>

```
[44]: import plotly.graph_objects as go
import pandas as pd
import os

# Assuming clf is your trained RandomForestClassifier and X is your feature_
↪matrix
importances = clf.feature_importances_
features = X.columns
```

```

# Create a DataFrame with features and their importances
df_importances = pd.DataFrame({
    'feature': features,
    'importance': importances
})

# Sort the DataFrame in descending order of importance
df_importances = df_importances.sort_values('importance', ascending=False)

# Create a Figure
fig = go.Figure()

# Add a trace for each feature
for _, row in df_importances.iterrows():
    fig.add_trace(go.Bar(x=[row['feature']], y=[row['importance']],
                        hovertemplate='Importance: %{y}<extra></extra>'))

# Set layout properties
fig.update_layout(
    title='Feature Importance',
    xaxis_title='Features',
    yaxis_title='Importance',
    showlegend=False,
    hovermode='x'
)

# Show the figure
fig.show()

```

```

[45]: avg_feature_importance = np.mean(feature_importance)
      avg_feature_importance

```

```

[45]: 0.11111111111111113

```

The meaning of having high value and low value of importance is as follows:

1. Importance is a measure of how much a feature contributes to the prediction accuracy of a machine learning model.
2. High value of importance means that the feature is very relevant and influential for the model's performance. Changing or removing the feature would significantly affect the model's accuracy.
3. Low value of importance means that the feature is not very relevant or influential for the model's performance. Changing or removing the feature would not significantly affect the model's accuracy.

In this chart, euribor3m, age, and nr.employed are greater than average importance score (based on heuristic approach). Thus, let's see how these three column characteristics based on target variable.

Reference:

https://scikit-learn.org/stable/auto_examples/ensemble/plot_forest_importances.html

```
[46]: # Convert X_resampled to DataFrame
X_resampled_df = pd.DataFrame(X_resampled, columns=X.columns)

# Convert y_resampled to DataFrame
y_resampled_df = pd.DataFrame(y_resampled, columns=['y'])

# Concatenate X_resampled_df and y_resampled_df
resampled_df = pd.concat([X_resampled_df, y_resampled_df], axis=1)

resampled_df.head()
```

```
[46]:   age  campaign  pdays  previous  emp.var.rate  cons.price.idx  \
0    56         1    999         0         1.1         93.994
1    57         1    999         0         1.1         93.994
2    37         1    999         0         1.1         93.994
3    40         1    999         0         1.1         93.994
4    56         1    999         0         1.1         93.994
```

```
   cons.conf.idx  euribor3m  nr.employed  y
0         -36.4      4.857      5191.0  0
1         -36.4      4.857      5191.0  0
2         -36.4      4.857      5191.0  0
3         -36.4      4.857      5191.0  0
4         -36.4      4.857      5191.0  0
```

```
[47]: resampled_df.groupby('y')['euribor3m'].describe()
```

```
[47]:   count      mean      std   min    25%    50%    75%   max
y
0  36548.0  3.811491  1.638187  0.634  1.405000  4.857000  4.962000  5.045
1  36548.0  2.120282  1.743075  0.634  0.847427  1.267079  4.330096  5.045
```

```
[48]: resampled_df.groupby('y')['age'].describe()
```

```
[48]:   count      mean      std   min    25%    50%    75%   max
y
0  36548.0  39.911185  9.898132  17.0   32.0   38.0   47.0  95.0
1  36548.0  40.640363  13.761825  17.0   30.0   37.0   49.0  98.0
```

```
[49]: resampled_df.groupby('y')['nr.employed'].describe()
```

```
[49]:   count      mean      std   min    25%    50%    75%   max
y
0  36548.0  5176.166600  64.571979  4963.6  5099.1  5195.8  5228.1  5228.1
1  36548.0  5095.003376  87.609914  4963.6  5017.5  5099.1  5191.0  5228.1
```



```
[50]: import plotly.express as px
fig = px.scatter_3d(resampled_df, x='euribor3m', y='age', z='nr.employed',
                    color='y')
fig.show()
```

Based on descriptive statistics and visualization, there is no significant difference between yes and no. Further investigation can be made using `class_weight` parameter in Random Forest and `train_test_split`. Thus, the temporary conclusion were in terms of bank marketing case, `euribor3m`, `age`, and `nr.employed` are the biggest factor that affecting whether customer will subscribe or not. Since it is assumed that both positive and negative response is equally importance, the use of F1 score is a good metrics to determine which model is better.